


## Efficient Point-to-Point Shortest Path Algorithms

Andrew V. Goldberg (Microsoft Research)  
Chris Harrelson (Google)  
Haim Kaplan (Tel Aviv University)  
Renato F. Werneck (Princeton University)

Northwest  
 $n = 1.6M$  vertices  
 $m = 3.8M$  arcs



3

## Shortest Paths

- Point-to-point shortest path problem (P2P):
  - Given:
    - \* directed graph with nonnegative arc lengths  $\ell(v, w)$ ;
    - \* source vertex  $s$ ;
    - \* target vertex  $t$ .
  - Goal: find shortest path from  $s$  to  $t$ .
- Our study:
  - Large road networks:
    - \* 330K (Bay Area) to 30M (North America) vertices.
  - Algorithms work in **two stages**:
    - \* preprocessing: may take hours, outputs linear amount of data;
    - \* query: should take milliseconds, uses the preprocessed data.

2

## Example Graph

## Obvious Algorithm

- Precompute all shortest paths and store distance matrix.
- Will not work on large graphs ( $n = 30M$ ).
  - $O(n^2)$  space:  $\sim 26$  PB.
  - $\tilde{O}(nm)$  time: years (single Dijkstra takes  $\sim 10s$ ).

(All times on a 2.4 GHz AMD Opteron with 16 GB of RAM.)


4

## Dijkstra's Algorithm

- Vertices processed in increasing order of distance:
  - maintains a **distance label**  $d(v)$  for each vertex:
    - \* upper bound on  $\text{dist}(s, v)$ ;
    - \* initially,  $d(s) = 0$  and  $d(v) = \infty$  for all other vertices.
  - In each iteration:
    - \* Pick unscanned vertex  $v$  with smallest  $d(\cdot)$  (use heap).
    - \* Scan  $v$ :
      - For each edge  $(v, w)$ , check if  $d(w) > d(v) + \ell(v, w)$ .
      - If it is, set  $d(w) \leftarrow d(v) + \ell(v, w)$ .
  - Stop when the target  $t$  is about to be scanned.
  - [Dijkstra'59, Dantzig'63].
- Intuition:
  - grow a ball around  $s$  and stop when  $t$  is scanned.

5

## Dijkstra's Algorithm




6

## Bidirectional Dijkstra's Algorithm

- Bidirectional Dijkstra's algorithm:
  - **forward** search from  $s$  with labels  $d_f$ :
    - \* performed on the **original graph**.
  - **reverse** search from  $t$  with labels  $d_r$ :
    - \* performed on the **reverse graph**;
    - \* same set of vertices, each arc  $(v, w)$  becomes  $(w, v)$ .
  - alternate in any way.
- Intuition: grow a ball around each end ( $s$  and  $t$ ) until they “meet”.

7

## Bidirectional Dijkstra's Algorithm



8

## Bidirectional Dijkstra's Algorithm

- Possible stopping criterion:
  - a vertex  $v$  is about to be scanned a second time:
    - \* once in each direction;
    - $v$  may not be on the shortest path.
- We must maintain the length  $\mu$  of the best path seen so far:
  - initially,  $\mu = \infty$ ;
  - when scanning an arc  $(v, w)$  in the forward search and  $w$  is scanned in the reverse search, update  $\mu$  if  $d_f(v) + \ell(v, w) + d_r(w) < \mu$ .
  - similar procedure if scanning an arc in the reverse search.

9

## Bidirectional Dijkstra's Algorithm

- Stronger stopping condition:
  - Let  $\text{top}_f$  and  $\text{top}_r$  be the top heap values (forward and reverse).
  - Stop when  $\text{top}_f + \text{top}_r \geq \mu$ .
  - Previous stopping criterion is a special case.
- Why does it work?
  - Suppose there exists an  $s-t$  path  $P$  with length less than  $\mu$ .
  - There must be an arc  $(v, w)$  on this path such that:
    - \*  $\text{dist}(s, v) < \text{top}_f$  and
    - \*  $\text{dist}(w, t) < \text{top}_r$ .
  - Both  $v$  and  $w$  have been scanned already.
  - When the second of these was scanned, it would have found the  $P$ .
    - \* Contradiction:  $P$  cannot exist.

10

## Part I: A\* Search

### A\* Search

- Define potential function  $\pi(v)$  and modify lengths:
  - $\ell_\pi(v, w) = \ell(v, w) - \pi(v) + \pi(w)$
  - $\ell_\pi(v, w)$ : reduced cost of arc  $(v, w)$ .
- All  $s-t$  paths change by same amount:  $\pi(t) - \pi(s)$ .
- A\* search:
  - Equivalent to Dijkstra on the modified graph:
    - \* correct if  $\ell_\pi(v, w) \geq 0$  ( $\pi$  feasible).
  - Vertices scanned in increasing order of  $k(v) = d(v) + \pi(v)$ :
    - \*  $\pi(v)$ : estimate on  $\text{dist}(v, t)$ ;
    - \*  $k(v)$ : estimated length of shortest  $s-t$  path through  $v$ .
  - If  $\pi(t) = 0$  and  $\pi$  feasible,  $\pi(v)$  is a lower bound on  $\text{dist}(v, t)$ .
- All we need are good feasible lower bounds (e.g., Euclidean).

11

12

## A\* Search

- Why is A\* equivalent to Dijkstra on the modified graph?
  - Dijkstra picks vertices with increasing (modified) distance from  $s$ :
    - \*  $\text{dist}_\pi(s, v) = \text{dist}(s, v) - \pi(s) + \pi(v)$
  - A\* search picks vertices with increasing key:
    - \*  $k(v) = \text{dist}(s, v) + \pi(v)$
  - $\pi(s)$  is constant: these orders are the same.
- Why is  $\pi(v)$  a lower bound on  $\text{dist}(v, t)$  when  $\pi$  is feasible and  $\pi(t) = 0$ ?
  - Take the shortest path from  $v$  to  $t$ .
  - Two ways of computing its reduced cost:
    1.  $\text{dist}(v, t) - \pi(v) + \pi(t) = \text{dist}(v, t) - \pi(v)$  (since  $\pi(t) = 0$ );
    2. sum of the reduced costs of all arcs:
      - \* must be nonnegative, since  $\pi$  is feasible.
  - Combining them:  $\text{dist}(v, t) - \pi(v) \geq 0 \Rightarrow \pi(v) \leq \text{dist}(v, t)$ .

13

## Bidirectional A\* Search

- Bidirectional search needs two potential functions:
  - $\pi_f(v)$ : estimate on  $\text{dist}(v, t)$ .
  - $\pi_r(v)$ : estimate on  $\text{dist}(s, v)$ .
- Reduced cost of arc  $(v, w)$ :
  - Forward:  $\ell_f(v, w) = \ell(v, w) - \pi_f(v) + \pi_f(w)$ .
  - Reverse:  $\ell_r(w, v) = \ell(v, w) - \pi_r(w) + \pi_r(v)$ .
    - \* the arc appears as  $(w, v)$  in the reverse graph.
- These values must be **consistent**:

$$\begin{aligned}\ell_f(v, w) &= \ell_r(w, v) \\ \ell(v, w) - \pi_f(v) + \pi_f(w) &= \ell(v, w) - \pi_r(w) + \pi_r(v) \\ \pi_f(w) + \pi_r(w) &= \pi_f(v) + \pi_r(v)\end{aligned}$$

- This must be true for all pairs  $(v, w)$ , i.e.,  $(\pi_f + \pi_r) = \text{constant}$ .

14

## Bidirectional A\* Search

- Must use consistent potential functions.
- In general, two arbitrary feasible functions  $\pi_f$  and  $\pi_r$  are **not** consistent.
- Their **average** is both feasible and consistent [Ikeda et al. 94]:
  - $p_f(v) = \frac{1}{2}(\pi_f(v) - \pi_r(v))$
  - $p_r(v) = \frac{1}{2}(\pi_r(v) - \pi_f(v)) = -p_f(v)$
- To make the algorithm more intuitive, we make:
  - $p_f(v) = \frac{1}{2}(\pi_f(v) - \pi_r(v)) + \frac{\pi_r(t)}{2}$
  - $p_r(v) = \frac{1}{2}(\pi_r(v) - \pi_f(v)) + \frac{\pi_f(s)}{2}$
  - Added terms are constant: functions still feasible and consistent.
  - When  $\pi_f$  and  $\pi_r$  are lower bounds,  $p_f(t) = 0$  and  $p_r(s) = 0$ .
- $p$  usually provides worse bounds than  $\pi$ :
  - still worth it in practice.

15

## Bidirectional A\* Search

- Standard bidirectional Dijkstra:
  - stop when  $\text{top}_f + \text{top}_r \geq \mu$ .
    - \*  $\text{top}_f$ : length of the path from  $s$  to top element of forward heap.
    - \*  $\text{top}_r$ : length of (reverse) path from  $t$  to top element of reverse heap.
    - \*  $\mu$ : best  $s-t$  path seen so far.
- Bidirectional A\* search: same, but on the **modified graph**:
  - Let  $v_f$  and  $v_r$  be the top elements in each heap;
  - Length of path  $s-v_f$  is  $d_f(v_f) + p_f(v_f) - p_f(s) = \text{top}_f - p_f(s)$ .
  - Length of reverse path  $t-v_r$  is  $d_r(v_r) + p_r(v_r) - p_r(t) = \text{top}_r - p_r(t)$ .
  - Stopping criterion:

$$[\text{top}_f - p_f(s)] + [\text{top}_r - p_r(t)] \geq [\mu - p_f(s) + p_f(t)]$$

- Simplifying and using  $p_f(t) = 0$ :

$$\text{top}_f + \text{top}_r \geq \mu + p_r(t).$$

16

## Lower Bounds

- Preprocessing:
  - select a constant number of **landmarks** (we use 16);
  - for each landmark, precompute distance to and from every vertex.

- Lower bounds use the **triangle inequality**:

$$\text{dist}(v, w) \geq \text{dist}(A, w) - \text{dist}(A, v)$$


$$\text{dist}(v, w) \geq \text{dist}(v, A) - \text{dist}(w, A)$$

$$\text{dist}(v, w) \geq \max\{\text{dist}(A, w) - \text{dist}(A, v), \text{dist}(v, A) - \text{dist}(w, A)\}$$

- A good landmark appears “before”  $v$  or “after”  $w$ .
- More than one landmark: pick maximum (still feasible).

17

## Query with Landmarks



18

## Experimental Results

- Northwest (1 649 045 vertices), 1000 random pairs:

METHOD	PREPROCESSING		QUERY		
	minutes	MB	avgscan	maxscan	ms
Bidirectional Dijkstra	—	28	518 723	1 197 607	340.74
Landmarks	4	132	16 276	150 389	12.05

- Vertices scanned:  $\sim 1\%$  on average,  $\sim 10\%$  on bad cases.

19

## Landmark Selection

- Landmark selection happens in two stages.
- Preprocessing:
  - Pick a small number of landmarks (we use **16**).
    - \* more landmarks: better queries, more space.
  - Store on disk distances to and from each landmark.
- Query ( $s$  and  $t$  known):
  - using all available landmarks is expensive;
  - pick a small subset (**2 to 6**) that is good for the search.

20

## Landmark Selection during Preprocessing

- Ultimate goal:
  - There should be a landmark “behind” every  $s-t$  pair.
  - Graphs are big, cannot evaluate this exactly: use heuristics.
    - \* All methods are quasi-linear.
- Algorithms:
  - Simple methods: `random`, `farthest`, `planar`;
  - `avoid`: adds landmarks “behind” regions not currently covered;
  - `maxcover`: avoid + local search:
    - \* goal: maximize #arcs with zero reduced cost.
- Best in practice is `maxcover`:
  - queries  $\sim 3$  times as fast as `random`;
  - preprocessing  $\sim 15$  times slower.

21

## Landmark Selection at Query Time

- Use only an **active** subset:
  - prefer landmarks that give the best lower bound on  $\text{dist}(s, t)$ .
- We use **dynamic selection**:
  - start with two landmarks (best forward + best reverse);
  - periodically check if a new landmark would help;
  - heaps rebuilt when landmarks added.
- Performance in practice:
  - picks only  $\sim 3$  landmarks;
  - fewer nodes visited than with any fixed number of landmarks.

22

## Part II: Reach

23

## Reaches


- Let  $v$  be a vertex on the **shortest** path  $P$  between  $s$  and  $t$ .
- **Reach** of  $v$  with respect to  $P$ :
 
$$\text{reach}(v, P) = \min\{\text{dist}(s, v), \text{dist}(t, v)\}$$
- Reach of  $v$  with respect to the whole graph:
 
$$\text{reach}(v) = \max_P \{\text{reach}(v, P)\},$$

over all shortest paths  $P$  that contain  $v$  [Gutman'04].
- Intuition:
  - vertices on highways have high reach;
  - vertices on local roads have low reach.

24

## Using Reaches

- Reaches can be used to prune the search during an  $s-t$  query.
- While scanning an edge  $(v, w)$ :
  - If  $\text{reach}(w) < \min\{d(s, v) + \ell(v, w), \text{LB}(w, t)\}$ , then  $w$  can be pruned.




- How do we obtain lower bounds?
  - Explicitly: Euclidean distances (Gutman's suggestion), landmarks.
  - Implicitly: make the search bidirectional.

25

## Implicit Bounds: Bidirectional Search


- Let  $R_t$  be the radius of the reverse search:
  - $R_t$  is the value of the top element in the reverse heap;
  - if  $w$  not labeled in the reverse direction, then  $d(w, t) \geq R_t$ .



- Pruning test:  $\text{reach}(w) < \min\{d(s, v) + \ell(v, w), R_t\}$ 
  - for best results, balance the forward and reverse searches by radius.

26

## Queries with Reaches



27

## Experimental Results

- Northwest (1 649 045 vertices), 1000 random pairs:

METHOD	PREPROCESSING		QUERY		
	minutes	MB	avgscan	maxscan	ms
Bidirectional Dijkstra	—	28	518 723	1 197 607	340.74
Landmarks	4	132	16 276	150 389	12.05
Reaches	1100	34	53 888	106 288	30.61

28

## Computing Reaches

- Trivial algorithm:
  - compute every  $s$ - $t$  path;
  - determine reach of each vertex on each path.
- Implementation:
  - Build shortest path tree  $T_r$  from each vertex  $r$ ;
  - Determine reach of each vertex  $v$  within the tree:
$$\text{reach}(v, T_r) = \min\{\text{depth}(v), \text{height}(v)\}$$
  - Take maximum over all  $r$ .
- Runs in  $\tilde{O}(nm)$  time:
  - overnight on Bay Area, years on North America.

29


## Computing Reaches

- Query still correct with **upper bounds** on reaches.
- We use iterative algorithm:
  1. find vertices with reach at most  $\epsilon$ ;  
– look only at **partial** shortest path trees (depth  $\sim 2\epsilon$ ).
  2. eliminate vertices with small reach;  
– if no vertices remain, stop;  
– otherwise, increase  $\epsilon$  and start another iteration.
- Use **penalties** to account for vertices already eliminated:  
– reaches no longer exact, but valid upper bounds
- Works well if many vertices are eliminated between iterations.

30

## Shortcuts


- Consider a sequence of vertices of degree two on the path below:



31

## Shortcuts


- Consider a sequence of vertices of degree two on the path below:  
– they all have high reach;



32

## Shortcuts


- Consider a sequence of vertices of degree two on the path below:
  - they all have high reach.
- Add a **shortcut**:
  - single edge bypassing a path (with same length).
  - assume ties are broken by taking path with fewer nodes.



33

## Shortcuts


- Consider a sequence of vertices of degree two on the path below:
  - they all have high reach.
- Add a **shortcut**:
  - single edge bypassing a path (with same length).
  - assume ties are broken by taking path with fewer nodes.



34

## Shortcuts


- Consider a sequence of vertices of degree two on the path below:
  - they all have high reach.
- Add a **shortcut**:
  - single edge bypassing a path (with same length).
  - assume ties are broken by taking path with fewer nodes.
- More shortcuts can be added recursively.



35

## Shortcuts

- Consider a sequence of vertices of degree two on the path below:
  - they all have high reach.
- Add a **shortcut**:
  - single edge bypassing a path (with same length).
  - assume ties are broken by taking path with fewer nodes.
- More shortcuts can be added recursively.




36

## Shortcuts

- Adding shortcuts during preprocessing:
  - speeds up queries (pruning more effective);
  - speeds up preprocessing (graph shrinks faster);
  - requires slightly more space (graph has more arcs).
- Shortcuts bypass vertices of degree two:
  - some have degree two in the original graph;
  - some acquire degree two as other vertices are eliminated.
- Sanders and Schultes [ESA'05]:
  - similar idea for hierarchy-based algorithm.

37

## Reaches with Shortcuts



38

## Experimental Results


- Northwest (1 649 045 vertices), 1000 random pairs:

METHOD	PREPROCESSING		QUERY		
	minutes	MB	avgscan	maxscan	ms
Bidirectional Dijkstra	—	28	518 723	1 197 607	340.74
Landmarks	4	132	16 276	150 389	12.05
Reaches	1100	34	53 888	106 288	30.61
Reaches+Shortcuts	17	100	2 804	5 877	2.39

39


## Reaches and Landmarks

- A\* search with landmarks can use reaches:
  - A\* gives the search a sense of direction.
  - Reaches make the search sparser.
- Landmarks have dual purpose:
  1. guide the search;
  2. provide lower bounds for reach-based pruning.



40

## Reaches and Landmarks (with Shortcuts)



## Experimental Results

- Northwest (1 649 045 vertices), 1000 random pairs:

METHOD	PREPROCESSING		QUERY		
	minutes	MB	avgscan	maxscan	ms
Bidirectional Dijkstra	—	28	518 723	1 197 607	340.74
Landmarks	4	132	16 276	150 389	12.05
Reaches	1100	34	53 888	106 288	30.61
Reaches+Shortcuts	17	100	2 804	5 877	2.39
Reaches+Shortcuts+Landmarks	21	204	367	1 513.82.39 874	171.957

## References

- Goldberg, Harrelson, and Werneck (in preparation):
  - Goldberg and Harrelson (SODA'05):
    - \* “ALT algorithm” ( $A^*$  search + Landmarks + Triangle inequality).
  - Goldberg and Werneck (Alenex'05):
    - \* improved preprocessing and queries;
    - \* Pocket PC implementation.
- Goldberg, Kaplan, and Werneck (2005):
  - reach with shortcuts +  $A^*$  search.

<http://www.cs.princeton.edu/~rwerneck/public.htm>